

```
In [1]: import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: data = pd.read_csv('Personalized_Diet_Recommendations.csv')
data.head()
```

	Patient_ID	Age	Gender	Height_cm	Weight_kg	BMI	Chronic_Disease	Blood_Pressure_Systolic	Blood_Pressure_Diastolic	Cholesterol_Level	...	Protein_Intake	Carbohydrate_Intake	Fat_Intake	Preferred_Cui
0	P00001	56	Other	163	66	24.84	None	175	75	219	...	105	179	143	Wes
1	P00002	69	Female	171	114	38.99	None	155	72	208	...	69	315	75	Mediterr
2	P00003	46	Female	172	119	40.22	None	137	101	171	...	183	103	148	Wes
3	P00004	32	Female	197	118	30.41	None	148	91	258	...	135	371	120	Wes
4	P00005	60	Female	156	109	44.79	Hypertension	160	109	260	...	167	298	48	In

5 rows × 30 columns

```
In [4]: null_counts = data.isnull().sum()
# Print the number of null values
print(f"{null_counts.sum()} null entries have been found in the dataset\n")
# Drop null values
data.dropna(inplace=True) # or df_data = df_data.dropna()
```

```
# Find and handle duplicates
duplicate_count = data.duplicated().sum()
# Print the number of duplicate entries
print(f"{duplicate_count} duplicate entries have been found in the dataset\n")
# Remove duplicates
data.drop_duplicates(inplace=True) # or df_data = df_data.drop_duplicates()
# Display relative message
print("All duplicates have been removed\n")
```

```
# Reset the indexes
data.reset_index(drop=True, inplace=True)
```

```
# Inspect the dataset for categorical columns
print("Categorical columns:", data.select_dtypes(include=['object']).columns.tolist(), '\n')
```

```
# Print the first 5 Lines
data.head()
```

0 null entries have been found in the dataset

0 duplicate entries have been found in the dataset

All duplicates have been removed

Categorical columns: ['Patient\_ID', 'Gender', 'Chronic\_Disease', 'Genetic\_Risk\_Factor', 'Allergies', 'Alcohol\_Consumption', 'Smoking\_Habit', 'Dietary\_Habits', 'Preferred\_Cuisine', 'Food\_Aversions', 'Recommended\_Meal\_Plan']

	Patient_ID	Age	Gender	Height_cm	Weight_kg	BMI	Chronic_Disease	Blood_Pressure_Systolic	Blood_Pressure_Diastolic	Cholesterol_Level	...	Protein_Intake	Carbohydrate_Intake	Fat_Intake	Preferred_Cui
0	P00001	56	Other	163	66	24.84	None	175	75	219	...	105	179	143	Wes
1	P00002	69	Female	171	114	38.99	None	155	72	208	...	69	315	75	Mediterr
2	P00003	46	Female	172	119	40.22	None	137	101	171	...	183	103	148	Wes
3	P00004	32	Female	197	118	30.41	None	148	91	258	...	135	371	120	Wes
4	P00005	60	Female	156	109	44.79	Hypertension	160	109	260	...	167	298	48	In

5 rows × 30 columns

```
In [5]: data.columns
```

```
Out[5]: Index(['Patient_ID', 'Age', 'Gender', 'Height_cm', 'Weight_kg', 'BMI',
   'Chronic_Disease', 'Blood_Pressure_Systolic',
   'Blood_Pressure_Diastolic', 'Cholesterol_Level', 'Blood_Sugar_Level',
   'Genetic_Risk_Factor', 'Allergies', 'Daily_Steps', 'Exercise_Frequency',
   'Sleep_Hours', 'Alcohol_Consumption', 'Smoking_Habit', 'Dietary_Habits',
   'Caloric_Intake', 'Protein_Intake', 'Carbohydrate_Intake', 'Fat_Intake',
   'Preferred_Cuisine', 'Food_Aversions', 'Recommended_Calories',
   'Recommended_Protein', 'Recommended_Carbs', 'Recommended_Fats',
   'Recommended_Meal_Plan'],
  dtype='object')
```

```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 30 columns):
 #   Column           Non-Null Count Dtype
 ---  -- 
 0   Patient_ID       5000 non-null  object
 1   Age              5000 non-null  int64
 2   Gender            5000 non-null  object
 3   Height_cm         5000 non-null  int64
 4   Weight_kg         5000 non-null  int64
 5   BMI               5000 non-null  float64
 6   Chronic_Disease   5000 non-null  object
 7   Blood_Pressure_Systolic 5000 non-null  int64
 8   Blood_Pressure_Diastolic 5000 non-null  int64
 9   Cholesterol_Level 5000 non-null  int64
 10  Blood_Sugar_Level 5000 non-null  int64
 11  Genetic_Risk_Factor 5000 non-null  object
 12  Allergies          5000 non-null  object
 13  Daily_Steps        5000 non-null  int64
 14  Exercise_Frequency 5000 non-null  int64
 15  Sleep_Hours        5000 non-null  float64
 16  Alcohol_Consumption 5000 non-null  object
 17  Smoking_Habit      5000 non-null  object
 18  Dietary_Habits     5000 non-null  object
 19  Caloric_Intake     5000 non-null  int64
 20  Protein_Intake     5000 non-null  int64
 21  Carbohydrate_Intake 5000 non-null  int64
 22  Fat_Intake          5000 non-null  int64
 23  Preferred_Cuisine   5000 non-null  object
 24  Food_Aversions      5000 non-null  object
 25  Recommended_Calories 5000 non-null  int64
 26  Recommended_Protein 5000 non-null  int64
 27  Recommended_Carbs    5000 non-null  int64
 28  Recommended_Fats     5000 non-null  int64
 29  Recommended_Meal_Plan 5000 non-null  object
dtypes: float64(2), int64(17), object(11)
memory usage: 1.1+ MB
```


```
In [7]: columns_to_drop = ['Patient_ID', 'Preferred_Cuisine', 'Food_Aversions', 'Recommended_Meal_Plan', 'Recommended_Calories',
'Recommended_Protein', 'Recommended_Carbs', 'Recommended_Fats']
```

```
In [8]: data.drop(columns=columns_to_drop, inplace=True)
```


```
In [9]: data.columns
```

```
Out[9]: Index(['Age', 'Gender', 'Height_cm', 'Weight_kg', 'BMI', 'Chronic_Disease',
'Blood_Pressure_Systolic', 'Blood_Pressure_Diastolic',
'Cholesterol_Level', 'Blood_Sugar_Level', 'Genetic_Risk_Factor',
'Allergies', 'Daily_Steps', 'Exercise_Frequency', 'Sleep_Hours',
'Alcohol_Consumption', 'Smoking_Habit', 'Dietary_Habits',
'Caloric_Intake', 'Protein_Intake', 'Carbohydrate_Intake',
'Fat_Intake'],
dtype='object')
```

```
In [10]: sns.countplot(x='Chronic_Disease', data=data, palette='hls')
plt.show()
```



```
In [11]: sns.countplot(x='Age', hue='Gender', data=data, palette='CMRmap')
plt.show()
```




```
In [12]: data['Chronic_Disease'].value_counts()
```


```
Out[12]: None      2043
Diabetes   1019
Heart Disease 749
Hypertension 693
Obesity    496
Name: Chronic_Disease, dtype: int64
```

```
In [13]: target = data['Chronic_Disease'].value_counts()
labels = ['None', 'Diabetes', 'Heart Disease', 'Hypertension', 'Obesity']
sizes = [target[0], target[1], target[2], target[3], target[4]]
```

```
# Create pie chart
plt.figure(figsize=(6, 6))
plt.pie(sizes, labels=labels, autopct='%.1f%%',
        startangle=90, explode =[0.1, 0, 0.2, 0.3, 0.4])
plt.title('Sample Outcome')
plt.axis('equal')
plt.tight_layout()
plt.show()
```



```
In [14]: plt.figure(figsize = (15,10))
sns.heatmap(data.corr(), annot = True, cmap="rainbow")
plt.title('Correlation Matrics')
plt.show()
```



```
In [15]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 22 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              5000 non-null    int64  
 1   Gender            5000 non-null    object  
 2   Height_cm         5000 non-null    int64  
 3   Weight_kg         5000 non-null    int64  
 4   BMI               5000 non-null    float64 
 5   Chronic_Disease  5000 non-null    object  
 6   Blood_Pressure_Systolic 5000 non-null    int64  
 7   Blood_Pressure_Diastolic 5000 non-null    int64  
 8   Cholesterol_Level 5000 non-null    int64  
 9   Blood_Sugar_Level 5000 non-null    int64  
 10  Genetic_Risk_Factor 5000 non-null    object  
 11  Allergies          5000 non-null    object  
 12  Daily_Steps        5000 non-null    int64  
 13  Exercise_Frequency 5000 non-null    int64  
 14  Sleep_Hours       5000 non-null    float64 
 15  Alcohol_Consumption 5000 non-null    object  
 16  Smoking_Habit      5000 non-null    object  
 17  Dietary_Habits    5000 non-null    object  
 18  Caloric_Intake    5000 non-null    int64  
 19  Protein_Intake    5000 non-null    int64  
 20  Carbohydrate_Intake 5000 non-null    int64  
 21  Fat_Intake         5000 non-null    int64  
dtypes: float64(2), int64(13), object(7)
memory usage: 859.5+ KB
```

In [16]: `data.columns`

```
Out[16]: Index(['Age', 'Gender', 'Height_cm', 'Weight_kg', 'BMI', 'Chronic_Disease',
 'Blood_Pressure_Systolic', 'Blood_Pressure_Diastolic',
 'Cholesterol_Level', 'Blood_Sugar_Level', 'Genetic_Risk_Factor',
 'Allergies', 'Daily_Steps', 'Exercise_Frequency', 'Sleep_Hours',
 'Alcohol_Consumption', 'Smoking_Habit', 'Dietary_Habits',
 'Caloric_Intake', 'Protein_Intake', 'Carbohydrate_Intake',
 'Fat_Intake'],  
dtype='object')
```

In [17]: `data = data[['Age', 'Gender', 'Height_cm', 'Weight_kg', 'BMI',
 'Blood_Pressure_Systolic', 'Blood_Pressure_Diastolic',
 'Cholesterol_Level', 'Blood_Sugar_Level', 'Genetic_Risk_Factor',
 'Allergies', 'Daily_Steps', 'Exercise_Frequency', 'Sleep_Hours',
 'Alcohol_Consumption', 'Smoking_Habit', 'Dietary_Habits',
 'Caloric_Intake', 'Protein_Intake', 'Carbohydrate_Intake',
 'Fat_Intake', 'Chronic_Disease']]`

In [18]: `data.to_csv('processed.csv')`

In [19]: `print("Categorical columns:", data.select_dtypes(include=['object']).columns.tolist(), '\n')`

```
Categorical columns: ['Gender', 'Genetic_Risk_Factor', 'Allergies', 'Alcohol_Consumption', 'Smoking_Habit', 'Dietary_Habits', 'Chronic_Disease']
```

In [20]: `# Import Label encoder
from sklearn import preprocessing

# LabelEncoder object knows
# how to understand word Labels.
label_encoder = preprocessing.LabelEncoder()

# Encode Labels in column 'species'.
data['Gender']= label_encoder.fit_transform(data['Gender'])
data['Genetic_Risk_Factor']= label_encoder.fit_transform(data['Genetic_Risk_Factor'])
data['Allergies']= label_encoder.fit_transform(data['Allergies'])
data['Alcohol_Consumption']= label_encoder.fit_transform(data['Alcohol_Consumption'])
data['Smoking_Habit']= label_encoder.fit_transform(data['Smoking_Habit'])
data['Dietary_Habits']= label_encoder.fit_transform(data['Dietary_Habits'])
data['Chronic_Disease']= label_encoder.fit_transform(data['Chronic_Disease'])`

In [21]: `ML_Model = []
accuracy = []
precision = []
recall = []
f1score = []


def storeResults(model, a, b, c, d):
 ML_Model.append(model)
 accuracy.append(round(a, 3))
 precision.append(round(b, 3))
 recall.append(round(c, 3))
 f1score.append(round(d, 3))`

In [22]: `X = data.drop(["Chronic_Disease"],axis =1)
y = data["Chronic_Disease"]`

In [23]: `from imblearn.over_sampling import RandomOverSampler
sm = RandomOverSampler(random_state=42)
X_res, y_res = sm.fit_resample(X, y)`

In [24]: `target = y_res.value_counts()
labels = ['None', 'Diabetes', 'Heart Disease', 'Hypertension', 'Obesity']
sizes = [target[0], target[1], target[2], target[3], target[4]]

# Create bar chart
plt.figure(figsize=(6, 4))
plt.bar(labels, sizes)
plt.title('Balanced Disease Sample Outcome')
plt.ylabel('Number of Samples')
plt.tight_layout()
plt.show()`



```
In [25]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_res)

In [26]: import os, joblib
os.makedirs('Models', exist_ok=True)
joblib.dump(scaler, 'Models/scaler.sav')

Out[26]: ['Models/scaler.sav']

In [27]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, ConfusionMatrixDisplay, classification_report

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_res, test_size=0.3)
```

## Random Forest

```
In [29]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(
    n_estimators=500,
    max_depth=100,
    n_jobs=-1,
    random_state=42
)
rf.fit(X_train, y_train)

y_pred = rf.predict(X_test)


rf_acc = accuracy_score(y_pred, y_test)
rf_prec = precision_score(y_pred, y_test, average='weighted')
rf_rec = recall_score(y_pred, y_test, average='weighted')
rf_f1 = f1_score(y_pred, y_test, average='macro')

In [30]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.82	0.69	0.75	608
1	0.94	0.86	0.90	594
2	0.96	0.89	0.92	626
3	0.61	0.81	0.69	628
4	0.98	0.96	0.97	609
accuracy			0.84	3065
macro avg	0.86	0.84	0.85	3065
weighted avg	0.86	0.84	0.85	3065

```
In [31]: cm = confusion_matrix(y_test, y_pred)

# Plot Confusion Matrix
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```



```
In [32]: storeResults('Random Forest', rf_acc, rf_prec, rf_rec, rf_f1)
```

## Decision Tree

```
In [33]: from sklearn.tree import DecisionTreeClassifier
dt = DecisionTreeClassifier()
```

```
    max_depth=1000,
    random_state=42
)

dt.fit(X_train, y_train)

y_pred = dt.predict(X_test)


dt_acc = accuracy_score(y_pred, y_test)
dt_prec = precision_score(y_pred, y_test, average='weighted')
dt_rec = recall_score(y_pred, y_test, average='weighted')
dt_f1 = f1_score(y_pred, y_test, average='macro')
```

In [34]: `print(classification_report(y_test, y_pred))`

	precision	recall	f1-score	support
0	0.75	0.75	0.75	608
1	0.80	0.89	0.84	594
2	0.81	0.91	0.86	626
3	0.64	0.41	0.50	628
4	0.84	0.96	0.89	609
accuracy			0.78	3065
macro avg	0.77	0.78	0.77	3065
weighted avg	0.77	0.78	0.77	3065

In [35]: `cm = confusion_matrix(y_test, y_pred)`

```
# Plot Confusion Matrix
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```

In [36]: `storeResults('DecisionTree', dt_acc, dt_prec, dt_rec, dt_f1)`

## Logistic Regression

```
In [37]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(max_iter=100, C=0.1, random_state=42)

lr.fit(X_train, y_train)
```


Out[37]: `LogisticRegression(C=0.1, random_state=42)`

```
In [38]: y_pred = lr.predict(X_test)

lr_acc = accuracy_score(y_pred, y_test)
lr_prec = precision_score(y_pred, y_test, average='weighted')
lr_rec = recall_score(y_pred, y_test, average='weighted')
lr_f1 = f1_score(y_pred, y_test, average='weighted')
```

In [39]: `cm = confusion_matrix(y_test, y_pred)`

```
# Plot Confusion Matrix
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```



```
In [40]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.25	0.30	0.27	608
1	0.24	0.30	0.27	594
2	0.25	0.17	0.20	626
3	0.27	0.18	0.21	628
4	0.23	0.28	0.25	609
accuracy			0.24	3065
macro avg	0.25	0.25	0.24	3065
weighted avg	0.25	0.24	0.24	3065

```
In [41]: storeResults('Logistic Regression', lr_acc, lr_prec, lr_rec, lr_f1)
```

## SVM

```
In [42]: from sklearn.kernel_approximation import RBFSampler
from sklearn.svm import LinearSVC
from sklearn.pipeline import make_pipeline

svm = make_pipeline(
    RBFSampler(gamma=0.1, n_components=500),
    LinearSVC(max_iter=100)
)
```


```
In [43]: svm.fit(X_train, y_train)
```

```
Out[43]: Pipeline(steps=[('rbfsampler', RBFSampler(gamma=0.1, n_components=500)),
 ('linearsvc', LinearSVC(max_iter=100))])
```

```
In [44]: y_pred = svm.predict(X_test)

svm_acc = accuracy_score(y_pred, y_test)
svm_prec = precision_score(y_pred, y_test, average='weighted')
svm_rec = recall_score(y_pred, y_test, average='weighted')
svm_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [45]: cm = confusion_matrix(y_test, y_pred)
```



```
In [46]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.35	0.34	0.34	608
1	0.43	0.47	0.45	594
2	0.43	0.45	0.44	626
3	0.26	0.17	0.20	628
4	0.42	0.53	0.47	609
accuracy			0.39	3065
macro avg	0.38	0.39	0.38	3065
weighted avg	0.38	0.39	0.38	3065

```
In [47]: storeResults('SVM', svm_acc, svm_prec, svm_rec, svm_f1)
```

## MLP

```
In [48]: from sklearn.neural_network import MLPClassifier

mlp = MLPClassifier(random_state=1, max_iter=300)
mlp.fit(X_train, y_train)
```


```
Out[48]: MLPClassifier(max_iter=300, random_state=1)
```

```
In [49]: y_pred = mlp.predict(X_test)

mlp_acc = accuracy_score(y_pred, y_test)
mlp_prec = precision_score(y_pred, y_test, average='weighted')
mlp_rec = recall_score(y_pred, y_test, average='weighted')
mlp_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [50]: cm = confusion_matrix(y_test, y_pred)
```





```
In [51]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.44	0.44	0.44	608
1	0.53	0.54	0.53	594
2	0.53	0.58	0.56	626
3	0.31	0.21	0.25	628
4	0.57	0.70	0.63	609
accuracy			0.49	3065
macro avg	0.48	0.49	0.48	3065
weighted avg	0.48	0.49	0.48	3065

```
In [52]: storeResults('MLP',mlp_acc,mlp_prec,mlp_rec,mlp_f1)
```

## XGBoost

```
In [53]: from xgboost import XGBClassifier


xgb = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss')
xgb.fit(X_train, y_train)

y_pred = xgb.predict(X_test)

xgb_acc = accuracy_score(y_pred, y_test)
xgb_prec = precision_score(y_pred, y_test, average='weighted')
xgb_rec = recall_score(y_pred, y_test, average='weighted')
xgb_f1 = f1_score(y_pred, y_test, average='weighted')
```

```
In [54]: cm = confusion_matrix(y_test, y_pred)
```

```
# Plot Confusion Matrix
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```



```
In [55]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.73	0.71	0.72	608
1	0.81	0.89	0.84	594
2	0.85	0.90	0.87	626
3	0.65	0.54	0.59	628
4	0.91	0.96	0.93	609
accuracy			0.80	3065
macro avg	0.79	0.80	0.79	3065
weighted avg	0.79	0.80	0.79	3065

```
In [56]: storeResults('XGBoost',xgb_acc,xgb_prec,xgb_rec,xgb_f1)
```

## Ensemble

```
In [57]: from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import AdaBoostClassifier, VotingClassifier, BaggingClassifier
```

```
In [58]: rf = RandomForestClassifier(random_state=42, n_jobs=-1)
```

```
param_grid_rf = {
    "n_estimators": [100, 200, 300],
    "max_depth": [None, 10, 20],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 4],
    "max_features": ["sqrt", "log2"]
}

grid_rf = GridSearchCV(rf, param_grid_rf, cv=3, scoring="accuracy", n_jobs=-1, verbose=2)
grid_rf.fit(X_train, y_train)
print("Best RF params:", grid_rf.best_params_)
print("Best RF score:", grid_rf.best_score_)
```

Fitting 3 folds for each of 162 candidates, totalling 486 fits  
 Best RF params: {'max\_depth': None, 'max\_features': 'sqrt', 'min\_samples\_leaf': 1, 'min\_samples\_split': 5, 'n\_estimators': 300}  
 Best RF score: 0.7050340733814556

```
In [59]: best_rf = grid_rf.best_estimator_
```

```
In [60]: brf = BaggingClassifier(
    base_estimator=best_rf,
    random_state=42,
    n_jobs=-1
)
```

```
In [61]: ext = VotingClassifier(
    estimators=[('RF', best_rf), ('BagRF', brf)],
    voting='soft'
)
```

```
In [62]: ext.fit(X_train, y_train)
```

```
Out[62]: VotingClassifier(estimators=[('RF',
    RandomForestClassifier(max_features='sqrt',
        min_samples_split=5,
        n_estimators=300,
        n_jobs=-1,
        random_state=42)),
    ('BagRF',
        BaggingClassifier(base_estimator=RandomForestClassifier(max_features='sqrt',
            min_samples_split=5,
            n_estimators=300,
            n_jobs=-1,
            random_state=42),
            n_jobs=-1, random_state=42))),
    voting='soft')
```

```
In [63]: y_pred = ext.predict(X_test)
```


```
In [64]: ext_acc = accuracy_score(y_pred, y_test)
ext_prec = precision_score(y_pred, y_test, average='weighted')
ext_rec = recall_score(y_pred, y_test, average='weighted')
ext_f1 = f1_score(y_pred, y_test, average='macro')
```

```
In [65]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.70	0.75	608
1	0.94	0.86	0.90	594
2	0.95	0.89	0.92	626
3	0.62	0.78	0.69	628
4	0.97	0.96	0.96	609
accuracy			0.84	3065
macro avg	0.85	0.84	0.84	3065
weighted avg	0.85	0.84	0.84	3065

```
In [66]: cm = confusion_matrix(y_test, y_pred)
```

```
# Plot Confusion Matrix
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```



```
In [67]: storeResults('Ensemble Model', ext_acc, ext_prec, ext_rec, ext_f1)
```

## Comparison

```
In [68]: result = pd.DataFrame({ 'ML Model' : ML_Model,
    'Accuracy' : accuracy,
    'Macro F1 score' : f1score,
    'Recall' : recall,
    'Precision': precision,
})
```

In [69]: result




	ML Model	Accuracy	Macro F1 score	Recall	Precision
0	Random Forest	0.840	0.846	0.840	0.841
1	DecisionTree	0.779	0.767	0.779	0.818
2	Logistic Regression	0.244	0.248	0.244	0.260
3	SVM	0.388	0.398	0.388	0.414
4	MLP	0.492	0.503	0.492	0.521
5	XGBoost	0.796	0.801	0.796	0.810
6	Ensemble Model	0.837	0.842	0.837	0.836


## Graph

In [71]: palettes = ['Set2', 'Paired', 'tab10', 'Dark2']

```
import matplotlib.pyplot as plt
import seaborn as sns

for metric, pal in zip(['Accuracy', 'Macro F1 score', 'Recall', 'Precision'], palettes):
    plt.figure(figsize=(8, 4))
    sns.barplot(data=result, x='ML Model', y=metric, palette=sns.color_palette(pal))
    plt.title(f'{metric} by Model')
    plt.xticks(rotation=45)
    plt.tight_layout()
    plt.show()
```





## Models

```
In [78]: import joblib
```

```
filename = 'Models/model.sav'  
joblib.dump(ext, filename)
```

```
Out[78]: ['Models/model.sav']
```

```
In [79]: filename = 'Models/model_rf.sav'  
joblib.dump(rf, filename)
```

```
Out[79]: ['Models/model_rf.sav']
```

```
In [ ]:
```